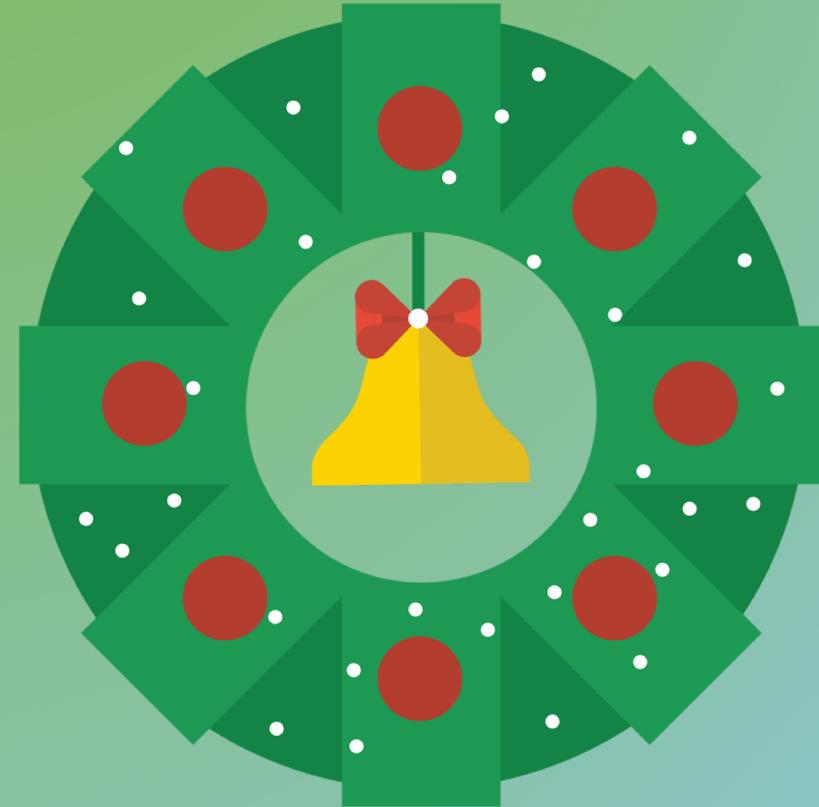


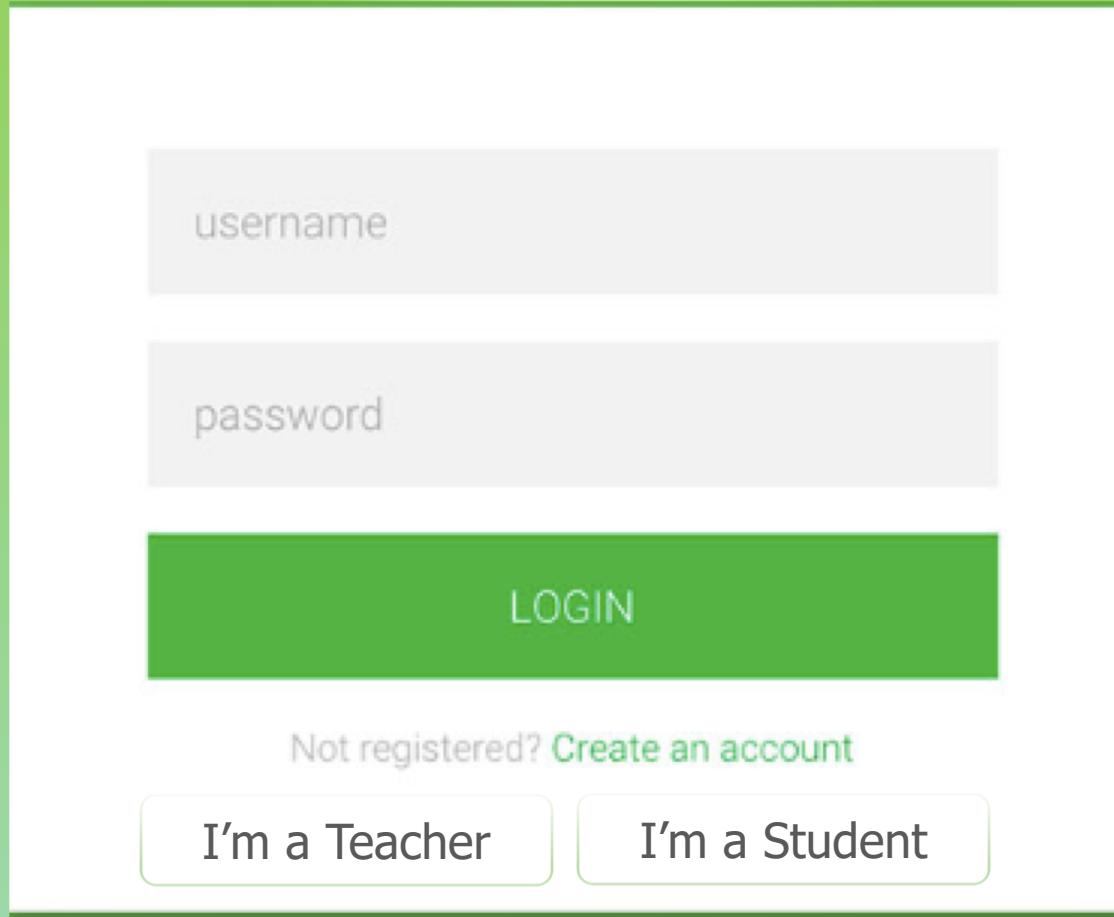
# Shared Notes

**Create Your Own Notes  
Enjoy Your Study  
Be happy every day  
Have Fun**

*Login / Register*



# Shared Notes



The image shows a login form with a light gray background and a dark green header and footer. The header contains the text "Shared Notes" in white. The footer has a solid green background with white text.

username

password

LOGIN

Not registered? [Create an account](#)

I'm a Teacher

I'm a Student

# Shared Notes



## Trace-based Just-in-Time Type Specialization for Dynamic Languages

Andreas Gal<sup>†,‡</sup>, Brendan Eich<sup>\*</sup>, Mike Shaver<sup>\*</sup>, David Anderson<sup>\*</sup>, David Mandelin<sup>\*</sup>,  
Mohammad R. Haghhighi<sup>§</sup>, Blake Kaplan<sup>¶</sup>, Graydon Hoare<sup>¶</sup>, Boris Zhuksky<sup>\*</sup>, Jason Orendorff<sup>\*</sup>,  
Jesse Ruderman<sup>\*</sup>, Edwin Smith<sup>||</sup>, Rick Reitmaier<sup>\*</sup>, Michael Bebenita<sup>\*</sup>, Mason Chang<sup>¶,‡</sup>, Michael Franz<sup>‡</sup>  
<sup>Mozilla Corporation\*</sup>  
{gal,brendan,shaver,danderson,mandelin,mrkecap,graydon,bsz,jroderoff,jruderman}@mozilla.com  
<sup>Adobe Corporation<sup>¶</sup></sup>  
{edwinsmith,creighton}@adobe.com  
<sup>Intel Corporation<sup>§</sup></sup>  
{mohammed.r.haghhighi}@intel.com  
<sup>University of California, Irvine<sup>¶</sup></sup>  
{mbebenita,mcchang,franz}@uci.edu

### Abstract

Dynamic languages such as JavaScript are more difficult to compile than statically typed ones. Since no concrete type information is available, traditional compilers need to emit generic code that can handle all possible type combinations at runtime. We present an alternative compilation technique for dynamically typed languages that allows us to generate machine code on the fly that is specialized for external dynamic types occurring on each path through the loop. Our method provides cheap *post-parsed* type specialization, and as elegant and efficient way of incrementally compiling (only) discovered alternative paths through nested loops. We have implemented a dynamic compiler framework based on our technique and we have measured speeds of 10x and more for certain benchmarks programs.

**Categories and Subject Descriptors** D.3.4 [Programming Languages]: Processors — Interpretive compilers; code generation.

**General Terms** Design, Experimentation, Measurement, Performance.

**Keywords** JavaScript, just-in-time compilation, trace trees.

### 1. Introduction

Dynamic languages such as JavaScript, Python, and Ruby, are popular since they are expressive, accessible to non-experts, and make deployment as easy as distributing a source file. They are used for small scripts as well as for complex applications. JavaScript, for example, is the de facto standard for client-side web programming

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for commercial advantage and that the copyright notice and the full page header appear on the first page. To copy otherwise, to republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee.

PLDI'09, June 15–30, 2009, Dublin, Ireland.  
Copyright © 2009 ACM 978-1-6053-6099-3/09/06...\$15.00.

and is used for the application logic of browser-based productivity applications such as Google Mail, Google Docs and Zimbra Collaboration Suite. In this domain, in order to provide a fluid user experience and enable a new generation of applications, visual interfaces are required that are generated on the fly.

Compilers for statically typed languages rely on type information to generate efficient machine code. In a dynamically typed programming language such as JavaScript, the types of expressions may vary at runtime. This means that the compiler can no longer easily transform operations into machine instructions that operate on one specific type. Without static type information, the compiler would need to generate specialized machine code for every possible type combination. While compilers for static type inference might be able to gather type information to generate optimized machine code, traditional static analysis is very expensive and hence not well suited for the highly interactive environment of a web browser.

We propose a trace-based compilation technique for dynamic languages that recoups speed of compilation with minimal performance of the generated machine code. Our system uses a staged-code execution approach: the system starts running JavaScript in a two-threaded bytecode interpreter. As the program runs, the system identifies hot (frequently executed) bytecode sequences, records them, and compiles them to fast native code. We can take a sequence of frames from a trace.

Unlike method-based dynamic compilers, our dynamic compiler operates at the granularity of individual loops. This design choice is based on the expectation that programs spend most of their time in hot loops. Even in dynamically typed languages, we expect hot loops to be mostly type stable, meaning that the types of variables involved in (12) For example, we want to keep pointers that start as integers remain integers for all iterations. When both of these expectations hold, a trace-based compiler can cover the program execution with a small number of type-specialized, efficiently compiled traces.

Each compiled trace covers one path through the program with one mapping of values to types. When the VM executes a compiled trace, it cannot guarantee that the same path will be followed or that the same types will occur in subsequent loop iterations.

Hence, recording and compiling a trace guarantees that the path and typing will be exactly as they were during recording for subsequent iterations of the loop.

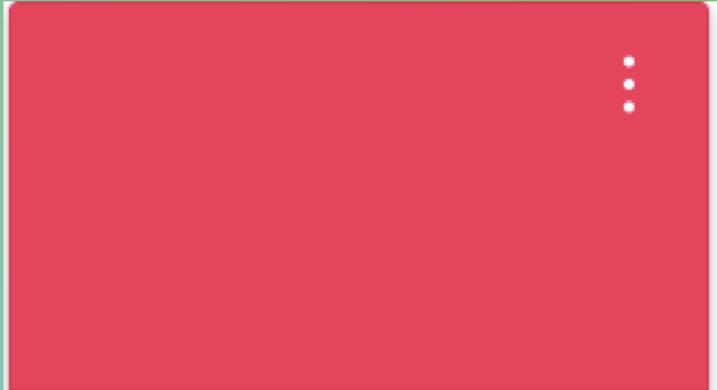
Every compiled trace contains all the guards (checks) required to validate the speculation. If one of the guards fails (if control flow is different, or a value of a different type is generated), the trace is discarded and a new one is recorded. The VM always traces starting at the exit to cover all new paths. In this way, the VM records a trace tree covering all the hot paths through the loops.

Nested loops can be difficult to optimize for tracing VMs. In a naive implementation, inner loops would become hot first, and the VM would start tracing them. When the inner loop exits, the VM would detect that a different branch was taken. The VM would

```
1 for (var i = 2; i < 100; ++i) {
2   if (!primes[i])
3     continue;
4   for (var k = i + 1; k < 100; k += i)
5     prime[k] = false;
6 }
```

Figure 1. Sample program: sieve of Eratosthenes, primes is initialized to an array of 100 false values on entry to this code snippet.

# Shared Notes



[Introduction to Information Security](#)

[18631](#)  
[Fall 2018](#)



[Introduction to Machine Learning...](#)

[10601-AC](#)  
[Fall 2018](#)



[Introduction to Machine Learning...](#)

[10601-AC](#)  
[Fall 2018](#)

[\*\*View Notes\*\*](#)

[\*\*View Notes\*\*](#)

[\*\*View Notes\*\*](#)

# Shared Notes

Edit

Edit

Edit

Bookmark

Bookmark

Bookmark

Create A New Post